

## Targets and target types

Please refer to the [simple\\_bev](#) code.

At [line 81](#), these are the data for one sample. (Note: not all of them are useful)

```
def run_model(model, loss_fn, d, device='cuda:0', sw=None):
    metrics = {}
    total_loss = torch.tensor(0.0, requires_grad=True).to(device)

    imgs, rots, trans, intrins, pts0, extra0, pts, extra, lrtlist_velo,
    vislist, tidlist, scorelist, seg_bev_g, valid_bev_g, center_bev_g,
    offset_bev_g, radar_data, egopose = d

    # shapes and types
    # -----
    # imgs          : torch.Size([1, 1, 4, 3, 448, 672]) torch.float32
    # rots          : torch.Size([1, 1, 4, 3, 3]) torch.float32
    # trans         : torch.Size([1, 1, 4, 3]) torch.float32
    # intrins       : torch.Size([1, 1, 4, 4, 4]) torch.float32
    # pts0         : torch.Size([1, 1, 3, 24247]) torch.float32
    # extra0       : torch.Size([1, 1, 3, 24247]) torch.float32
    # pts          : torch.Size([1, 1, 3, 30000]) torch.float32
    # extra        : torch.Size([1, 1, 3, 30000]) torch.float32
    # lrtlist_velo : torch.Size([1, 1, 150, 19]) torch.float32
    # vislist      : torch.Size([1, 1, 150]) torch.float32
    # tidlist      : torch.Size([1, 1, 150]) torch.int64
    # scorelist    : torch.Size([1, 1, 150]) torch.float32
    # seg_bev_g    : torch.Size([1, 1, 1, 400, 400]) torch.float32
    # valid_bev_g  : torch.Size([1, 1, 1, 400, 400]) torch.float32
    # center_bev_g: torch.Size([1, 1, 1, 400, 400]) torch.float32
    # offset_bev_g: torch.Size([1, 1, 2, 400, 400]) torch.float32
    # radar_data   : torch.Size([1, 1, 3, 700]) torch.float32
    # egopose      : torch.Size([1, 1, 4, 4]) torch.float32
```

According to [line 83](#), the idimensions of the `imgs` are `B0,T,S,C,H,W` where:

- `B`: batch
- `T`: time (useless)
- `S`: sequence (4 because there are four cameras)
- `C`: channel
- `H`: image height
- `W`: image width

At [lines 179-184](#), these are the outputs from its model, which are supposed to be the inputs of the next model.

```

_, feat_bev_e, seg_bev_e, center_bev_e, offset_bev_e = model(
    rgb_camXs=rgb_camXs,
    pix_T_cams=pix_T_cams,
    cam0_T_camXs=cam0_T_camXs,
    vox_util=vox_util,
    rad_occ_mem0=in_occ_mem0)

# shapes and types
# -----
# feat_bev_e   : torch.Size([1, 128, 400, 400]) torch.float32
# seg_bev_e    : torch.Size([1, 1, 400, 400]) torch.float32
# center_bev_e: torch.Size([1, 1, 400, 400]) torch.float32
# offset_bev_e: torch.Size([1, 2, 400, 400]) torch.float32

```

When you ignore the dimension that has size 1, the shapes of `seg_bev_e` and `seg_bev_g` match, as they are the prediction and ground truth, respectively. This is also true for `center_bev_e` and `offset_bev_e`.

The model is called `Segnet` and its heads that create those outputs can be found in the `Decoder`.

## Setup and installation

Install requirements and download the [repo](#) on GrootCompute

```

conda create -n bev python=3.10
conda activate bev
conda install pytorch==1.12.1 torchvision==0.13.1 cudatoolkit=11.3 -c
pytorch
pip install -U pip
pip install fire efficientnet_pytorch tensorboardX scikit-image pandas
nuscenes-devkit
pip install laspy[laszip]

git clone git@github.com:GrootCompute/simple_bev.git
git checkout fork # use this branch to process carla data

```

## Training and visualization

To run, first put `postprocessed.zip` at the base directory (i.e. one level up from `simple_bev`) and unzip it there. Then, enter `simple_bev` and call `train_carla.py`:

```

cd simple_bev/
python train_carla.py \
    --exp_name='rgb_mine' \
    --max_iters=501 \
    --log_freq=100 \
    --save_freq=100 \
    --batch_size=1 \
    --nworkers=1 \

```

```
--grad_acc=5 \  
--use_scheduler=True \  
--data_dir='../postprocessed' \  
--log_dir='logs_postprocessed' \  
--ckpt_dir='checkpoints' \  
--device_ids=[0]
```

Once training is done, the model can be found in the subdirectory `checkpoints/1x5_3e-4s_rgb_mine_<time>` (where `<time>` looks something like `17:34:03`). To run the visualization, enter `simple_bev` and call `vis_carla.py` and substitute in the subdirectory name at the line `--init_dir='checkpoints/...'`.

```
cd simple_bev/  
python vis_carla.py \  
    --max_iters=1 \  
    --nworkers=1 \  
    --data_dir='../postprocessed' \  
    --log_dir='logs_postprocessed' \  
    --init_dir='checkpoints/1x5_3e-4s_rgb_mine_17:34:03' \  
    --device_ids=[0]
```

## Other things

It seems like the segmentation maps generated by `simple_bev` are not usable. There are a lot of confusions with the coordinate systems and transformation matrices (in `nuscenesdataset_carla.py`). I don't think these can be fixed easily.


---

## Targets and target types (nuScenes)






I ran `train_nuscenes.py` for 1000 iterations using **nuScenes data** to obtain a quick model. Then I ran `vis_nuscenes.py` with that model for 5 "timesteps" and these are the outputs. Only `seg_bev_e` (shape: `(1, 1, 200, 200)`) from the outputs of the model is being used for visualization in the script.





```
_, feat_bev_e, seg_bev_e, center_bev_e, offset_bev_e = model(  
    rgb_camXs=rgb_camXs,  
    pix_T_cams=pix_T_cams,  
    cam0_T_camXs=cam0_T_camXs,  
    vox_util=vox_util,  
    rad_occ_mem0=in_occ_mem0)
```

Ground truth (i.e. `seg_bev_g`)





Segnet output (i.e. `seg_bev_e`)




Ground truth (i.e. `seg_bev_g`)Segnet output (i.e. `seg_bev_e`)

I dumped `center_bev_e` and `offset_bev_e` into a JSON file and inspected the outputs. For `center_bev_e` (shape: `(1, 1, 200, 200)`), it looks like this. Only the first timestep is shown:

Ground truth (i.e. `center_bev_g`)Segnet output (i.e. `center_bev_e`)

For `offset_bev_e` (shape: `(1, 2, 200, 200)`), it looks like this. Note: they are shown as x-offset and y-offset separately:



Mean, min, max of the arrays:

```
>>> center_bev_g.mean(), center_bev_g.min(), center_bev_g.max()
(0.021626132781710477, 0.0, 1.0)
>>> center_bev_e.mean(), center_bev_e.min(), center_bev_e.max()
(0.13331910948131118, 0.0012023173039779067, 0.8007904291152954)
>>> offset_bev_g.mean(), offset_bev_g.min(), offset_bev_g.max()
(-0.02825, -24.0, 22.0)
>>> offset_bev_e.mean(), offset_bev_e.min(), offset_bev_e.max()
(-0.4612532062917482, -13.117025375366211, 13.411355018615723)
```

## Loss functions

The outputs of the model are named `seg_bev_e` (segmentation map), `center_bev_e` (centerness), and `offset_bev_e` (offset). Based on my understanding, `seg_bev_e` is the main output that can be used for downstream tasks; while `center_bev_e` and `offset_bev_e` are auxiliary outputs that are meant to optimize the training of the model. This is an excerpt from [their paper](#):

acting as the segmentation task head. Following FIERY [21], we complement the segmentation head with auxiliary task heads for predicting centerness and offset, which serve to regularize the model. The offset head produces a vector field where, within each object mask, each vector points to the center of that object. We train the segmentation head with a cross-entropy `loss`, and supervise the centerness and offset fields with an L1 `loss`. We use an uncertainty-based learnable

At [lines 186-189](#), these are the loss functions applied to `seg_bev_e`, `center_bev_e`, and `offset_bev_e`.

```
ce_loss = loss_fn(seg_bev_e, seg_bev_g, valid_bev_g) # this is like
torch.nn.BCEWithLogitsLoss, but using a 'valid' mask to remove invalid
elements and calculate the "masked mean". Note that they set pos_weight to
2.13.
center_loss = balanced_mse_loss(center_bev_e, center_bev_g) # this is
like torch.nn.MSELoss, but can be supplied with a 'valid' mask, which is
not supplied here
offset_loss = torch.abs(offset_bev_e - offset_bev_g).sum(dim=1,
keepdim=True) # this is like torch.nn.L1Loss
offset_loss = utils.basic.reduce_masked_mean(offset_loss,
seg_bev_g * valid_bev_g) # a 'valid' mask is used here to remove invalid
elements and calculate the "masked mean"
```

*Prompt:* Try to understand the intentions behind the choices of loss functions.

- Generally speaking, for a feature map or segmentation map, where the output is a grid, and each cell contains a probability in the range of 0-1, the natural choice is binary cross-entropy ( `torch.nn.BCEWithLogitsLoss` ) or multi-class cross-entropy ( `torch.nn.CrossEntropyLoss` ). For deeper reasons, I refer you to [Logistic regression](#).
- Meanwhile, if the output represents a measurement (with arbitrary range) that has linear response (e.g. distance), the natural choices are L1 loss and L2 loss. If the measurement is non-linear, typically you want to linearized it first (e.g. if the output is exponential, you'd use  $\log(x)$  instead of  $x$ ). See [Linear regression](#).
  - L2 loss is very common but is sensitive to outliers (loss increases as  $x^2$  at large  $x$ ); L1 loss avoids that, but its gradient has a discontinuity at 0.
- For area-based outputs, people use IoU-related loss functions.
- And there are many other loss functions, e.g. focal loss which is optimized for cases where the two classes are highly imbalanced.
- Anyway, in the case of `simple_bev`, their loss functions are basically binary cross-entropy and L1/L2 losses (see above).